

DAT305

November 25, 2024

```
[1]: import sys,os
import pandas as pd
pd.options.mode.chained_assignment = None
import matplotlib.pyplot as plt
from matplotlib.patches import ConnectionPatch
import seaborn as sns
import numpy as np
import sklearn
import string
import re
import nltk
import tensorflow as tf
from collections import Counter
from tensorflow import keras
from sklearn.feature_extraction.text import TfidfVectorizer
from keras.utils import pad_sequences,to_categorical
from sklearn.feature_selection import SelectKBest,chi2
from sklearn.utils.class_weight import compute_class_weight
from sklearn.metrics import accuracy_score,f1_score,roc_auc_score,confusion_matrix,precision_score,recall_score,classif
from datetime import datetime
nltk.download('punkt')
nltk.download('stopwords')
from nltk.corpus import stopwords
from nltk.tokenize import word_tokenize
from nltk.stem import PorterStemmer
from nltk.tokenize import TweetTokenizer
from google.colab import files,drive
from wordcloud import WordCloud, STOPWORDS
from sklearn.svm import LinearSVC
from sklearn.model_selection import GridSearchCV
import warnings
warnings.filterwarnings(action="ignore", message="^internal gelsd")
print("Running Panda Version:"+pd.__version__)
print("Running TensorFlow Version:"+ tf.__version__)
#print("Running Keras API Version:"+ keras.__version__)
```

```
print("Running Python {0}.{1}".format(sys.version_info[:2][0],sys.version_info[:2][1]))
```

```
[nltk_data] Downloading package punkt to /root/nltk_data...
[nltk_data]   Unzipping tokenizers/punkt.zip.
[nltk_data] Downloading package stopwords to /root/nltk_data...
[nltk_data]   Unzipping corpora/stopwords.zip.
```

```
Running Panda Version:2.2.2
Running TensorFlow Version:2.17.1
Running Python 3.10
```

```
[2]: seed = 0
tf.keras.utils.set_random_seed(seed)
```

```
[3]: uploaded = files.upload()

<IPython.core.display.HTML object>

Saving Tweets.csv to Tweets.csv
```

1 *Exploratory Data Analysis*

```
[4]: dataset = pd.read_csv("Tweets.csv",na_values=['NA'], low_memory=False)
```

1.0.1 Dataset shapes

```
[5]: print('Dataset structure: rows =',dataset.shape[0], ' - columns =',dataset.
      ↪shape[1])
```

```
Dataset structure: rows = 14640 - columns = 15
```

Some random rows

```
[6]: dataset.sample(3)
```

```
[6]:
```

	tweet_id	airline_sentiment	airline_sentiment_confidence	\
13983	569682010270101504	negative	0.6163	
14484	569608307184242688	negative	0.7039	
6403	567879304593408001	negative	1.0000	

	negativereason	negativereason_confidence	airline	\
13983	Late Flight	0.6163	American	
14484	Bad Flight	0.3587	American	
6403	Cancelled Flight	1.0000	Southwest	

	airline_sentiment_gold	name	negativereason_gold	retweet_count	\
13983	NaN	zsalim03	NaN	0	
14484	NaN	sa_craig	NaN	0	
6403	NaN	DanaChristos	NaN	1	

		text	tweet_coord	\
13983	@AmericanAir	In car gng to DFW. Pulled over 1h...	NaN	
14484	@AmericanAir	after all, the plane didn't land ...	NaN	
6403	@SouthwestAir	can't believe how many paying cu...	NaN	

	tweet_created	tweet_location	\
13983	2015-02-22 18:15:50 -0800	Texas	
14484	2015-02-22 13:22:57 -0800	College Station, TX	
6403	2015-02-17 18:52:31 -0800	CT	

	user_timezone
13983	Central Time (US & Canada)
14484	Central Time (US & Canada)
6403	Eastern Time (US & Canada)

1.0.2 Descriptive statistics for the dataset

```
[7]: print('Dataset Features types:')
dataset.dtypes
```

Dataset Features types:

```
[7]: tweet_id          int64
airline_sentiment      object
airline_sentiment_confidence float64
negativereason         object
negativereason_confidence float64
airline                object
airline_sentiment_gold  object
name                   object
negativereason_gold    object
retweet_count          int64
text                   object
tweet_coord            object
tweet_created          object
tweet_location         object
user_timezone          object
dtype: object
```

```
[8]: print("List of names of columns:\n")
print('-'*40)
dataset.columns.tolist()
```

List of names of columns:

```
[8]: ['tweet_id',
      'airline_sentiment',
      'airline_sentiment_confidence',
      'negativereason',
      'negativereason_confidence',
      'airline',
      'airline_sentiment_gold',
      'name',
      'negativereason_gold',
      'retweet_count',
      'text',
      'tweet_coord',
      'tweet_created',
      'tweet_location',
      'user_timezone']
```

```
[9]: print('Descriptive Statistics for numeric features on Dataset')
      dataset.describe(include=np.number).T
```

Descriptive Statistics for numeric features on Dataset

```
[9]:
```

	count	mean	std \
tweet_id	14640.0	5.692184e+17	7.791112e+14
airline_sentiment_confidence	14640.0	9.001689e-01	1.628300e-01
negativereason_confidence	10522.0	6.382983e-01	3.304398e-01
retweet_count	14640.0	8.265027e-02	7.457782e-01

	min	25%	50% \
tweet_id	5.675883e+17	5.685592e+17	5.694779e+17
airline_sentiment_confidence	3.350000e-01	6.923000e-01	1.000000e+00
negativereason_confidence	0.000000e+00	3.606000e-01	6.706000e-01
retweet_count	0.000000e+00	0.000000e+00	0.000000e+00

	75%	max
tweet_id	5.698905e+17	5.703106e+17
airline_sentiment_confidence	1.000000e+00	1.000000e+00
negativereason_confidence	1.000000e+00	1.000000e+00
retweet_count	0.000000e+00	4.400000e+01

```
[10]: print('Range for numeric features on Dataset')
      print('-'*40)
      dataset.max(numeric_only=True) - dataset.min(numeric_only=True)
```

Range for numeric features on Dataset

```
[10]: tweet_id          2.722322e+15
      airline_sentiment_confidence  6.650000e-01
```

```
negativereason_confidence      1.000000e+00
retweet_count                  4.400000e+01
dtype: float64
```

```
[11]: dataset.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 14640 entries, 0 to 14639
Data columns (total 15 columns):
#   Column                                Non-Null Count  Dtype
---  -
0   tweet_id                             14640 non-null  int64
1   airline_sentiment                    14640 non-null  object
2   airline_sentiment_confidence         14640 non-null  float64
3   negativereason                       9178 non-null   object
4   negativereason_confidence            10522 non-null  float64
5   airline                              14640 non-null  object
6   airline_sentiment_gold                40 non-null     object
7   name                                 14640 non-null  object
8   negativereason_gold                  32 non-null     object
9   retweet_count                        14640 non-null  int64
10  text                                 14640 non-null  object
11  tweet_coord                           1019 non-null   object
12  tweet_created                         14640 non-null  object
13  tweet_location                       9907 non-null   object
14  user_timezone                         9820 non-null   object
dtypes: float64(2), int64(2), object(11)
memory usage: 1.7+ MB
```

```
[12]: print('Descriptive Statistics for categorical features on Dataset')
dataset.describe(include='O').T
```

Descriptive Statistics for categorical features on Dataset

```
[12]:
```

	count	unique	top	freq
airline_sentiment	14640	3	negative	9178
negativereason	9178	10	Customer Service Issue	2910
airline	14640	6	United	3822
airline_sentiment_gold	40	3	negative	32
name	14640	7701	JetBlueNews	63
negativereason_gold	32	13	Customer Service Issue	12
text	14640	14427	@united thanks	6
tweet_coord	1019	832	[0.0, 0.0]	164
tweet_created	14640	14247	2015-02-24 09:54:34 -0800	5
tweet_location	9907	3081	Boston, MA	157
user_timezone	9820	85	Eastern Time (US & Canada)	3744

```
[13]: print('Number of classes: ',len(dataset.airline_sentiment.unique().tolist()))
```

Number of classes: 3

```
[14]: print('Name of classes (y): ',dataset.airline_sentiment.unique().tolist())
```

Name of classes (y): ['neutral', 'positive', 'negative']

```
[15]: print('Check for duplicates?',dataset.duplicated().any())
print('-'*70)
print('Sum of duplicated rows :',dataset.duplicated().sum())
```

Check for duplicates? True

Sum of duplicated rows : 36

```
[16]: print('Check for missing values ',dataset.isnull().any().any())
print('-'*40)
print('Sum of Missing values accross columns\n',dataset.isnull().sum())
print('-'*40)
print("Sum of missing values",sum(dataset.isnull().sum()))
```

Check for missing values True

Sum of Missing values accross columns

tweet_id	0
airline_sentiment	0
airline_sentiment_confidence	0
negativereason	5462
negativereason_confidence	4118
airline	0
airline_sentiment_gold	14600
name	0
negativereason_gold	14608
retweet_count	0
text	0
tweet_coord	13621
tweet_created	0
tweet_location	4733
user_timezone	4820
dtype: int64	

Sum of missing values 61962

```
[17]: missing_data = dataset[dataset.isnull().any(axis=1)]
missing_data.head(3)
```

```
[17]:      tweet_id  airline_sentiment  airline_sentiment_confidence  \
0  570306133677760513             neutral             1.0000
1  570301130888122368             positive             0.3486
2  570301083672813571             neutral             0.6837
```

	negativereason	negativereason_confidence	airline	\
0	NaN	NaN	Virgin America	
1	NaN	0.0	Virgin America	
2	NaN	NaN	Virgin America	

	airline_sentiment_gold	name	negativereason_gold	retweet_count	\
0	NaN	cairdin	NaN	0	
1	NaN	jnardino	NaN	0	
2	NaN	yvonnalynn	NaN	0	

	text	tweet_coord	\
0	@VirginAmerica What @dhepburn said.	NaN	
1	@VirginAmerica plus you've added commercials t...	NaN	
2	@VirginAmerica I didn't today... Must mean I n...	NaN	

	tweet_created	tweet_location	user_timezone
0	2015-02-24 11:35:52 -0800	NaN	Eastern Time (US & Canada)
1	2015-02-24 11:15:59 -0800	NaN	Pacific Time (US & Canada)
2	2015-02-24 11:15:48 -0800	Lets Play	Central Time (US & Canada)

```
[18]: dataset.airline_sentiment.value_counts()
```

```
[18]: airline_sentiment
negative    9178
neutral     3099
positive    2363
Name: count, dtype: int64
```

Data imbalance, more negative and neutral sentiment than positive.

```
[19]: list_of_airlines =dataset.airline.unique().tolist()
print('list of airlines: ',list_of_airlines)
```

```
list of airlines:  ['Virgin America', 'United', 'Southwest', 'Delta', 'US
Airways', 'American']
```

```
[20]: print("Time of first tweet in the dataset:",dataset.tweet_created.min())
print('-'*65)
print("Time of last tweet in the dataset:",dataset.tweet_created.max())
```

```
Time of first tweet in the dataset: 2015-02-16 23:36:05 -0800
```

```
-----
```

```
Time of last tweet in the dataset: 2015-02-24 11:53:37 -0800
```

```
[21]: print("Airlines with tweet count\n", dataset.airline.value_counts())
```

```
Airlines with tweet count
airline
United          3822
```

```

US Airways      2913
American        2759
Southwest       2420
Delta           2222
Virgin America   504
Name: count, dtype: int64

```

```

[22]: print("Tweets frequencies grouped by sentiments for the airlines")
      print('-'*60)
      airlines_sentiments_groups = dataset.groupby("airline",
      ↪group_keys=True)[['airline_sentiment']].value_counts()
      airlines_sentiments_groups

```

Tweets frequencies grouped by sentiments for the airlines

```

[22]: airline      airline_sentiment
      American      negative      1960
              neutral      463
              positive      336
      Delta        negative      955
              neutral      723
              positive      544
      Southwest    negative     1186
              neutral      664
              positive      570
      US Airways   negative     2263
              neutral      381
              positive      269
      United       negative     2633
              neutral      697
              positive      492
      Virgin America negative      181
              neutral      171
              positive      152
      Name: count, dtype: int64

```

```

[23]: print("Maximum tweet confidence",dataset.airline_sentiment_confidence.max())
      print('-'*30)
      print("Minimum tweet confidence",dataset.airline_sentiment_confidence.min())

```

Maximum tweet confidence 1.0

Minimum tweet confidence 0.335

```

[24]: print(dataset.airline_sentiment_confidence.quantile([0,0.25,0.50,0.75,1]))

```

```

0.00    0.3350
0.25    0.6923

```



```
0.50    1.0000
0.75    1.0000
1.00    1.0000
```

```
Name: airline_sentiment_confidence, dtype: float64
```

#####25 percent quantile for sentiment confidence is 0.65, which means that 25 percent of the dataset values for this measure is less than 0.65.

```
[25]: dataset['negativereason'] = dataset['negativereason'].fillna('N/A')
```

```
[26]: dataset[dataset['negativereason'] != 'N/A'].groupby("airline",
↳group_keys=True)[['negativereason']].value_counts()
```

```
[26]: airline      negativereason
American      Customer Service Issue      768
              Late Flight                  249
              Cancelled Flight             246
              Can't Tell                   198
              Lost Luggage                 149
              Flight Booking Problems      130
              Bad Flight                   87
              Flight Attendant Complaints  87
              longlines                    34
              Damaged Luggage              12
Delta          Late Flight                  269
              Customer Service Issue      199
              Can't Tell                   186
              Bad Flight                   64
              Flight Attendant Complaints  60
              Lost Luggage                 57
              Cancelled Flight             51
              Flight Booking Problems      44
              longlines                    14
              Damaged Luggage              11
Southwest      Customer Service Issue      391
              Cancelled Flight             162
              Can't Tell                   159
              Late Flight                  152
              Bad Flight                   90
              Lost Luggage                 90
              Flight Booking Problems      61
              Flight Attendant Complaints  38
              longlines                    29
              Damaged Luggage              14
US Airways     Customer Service Issue      811
              Late Flight                  453
              Can't Tell                   246
              Cancelled Flight             189
```

	Lost Luggage	154
	Flight Attendant Complaints	123
	Flight Booking Problems	122
	Bad Flight	104
	longlines	50
	Damaged Luggage	11
United	Customer Service Issue	681
	Late Flight	525
	Can't Tell	379
	Lost Luggage	269
	Bad Flight	216
	Cancelled Flight	181
	Flight Attendant Complaints	168
	Flight Booking Problems	144
	longlines	48
	Damaged Luggage	22
Virgin America	Customer Service Issue	60
	Flight Booking Problems	28
	Can't Tell	22
	Bad Flight	19
	Cancelled Flight	18
	Late Flight	17
	Flight Attendant Complaints	5
	Lost Luggage	5
	Damaged Luggage	4
	longlines	3

Name: count, dtype: int64

Highest retweet

```
[27]: dataset[['text', 'airline', 'name', 'airline_sentiment']].
      ↪loc[dataset['retweet_count'].max()]
```

```
[27]: text                @VirginAmerica are flights leaving Dallas for ...
      airline                Virgin America
      name                papamurat
      airline_sentiment    neutral
      Name: 44, dtype: object
```

```
[27]:
```

1.1 Data visualization (Multivariate analysis).

```
[28]: target_classes, tweet_freq = np.unique(dataset.
      ↪airline_sentiment, return_counts=True)
      print(target_classes, tweet_freq)
```

```
['negative' 'neutral' 'positive'] [9178 3099 2363]
```

```

[29]: def func(pct, allvals)->str:
        absolute = int(np.round(pct/100.*np.sum(allvals)))
        return f"{pct:.1f}%\n ({absolute:d})"

[30]: # make figure and assign axis objects
fig, (ax1,ax2,ax3) = plt.subplots(1, 3, figsize=(20, 8))
fig.subplots_adjust(wspace=0)

# pie chart parameters
overall_ratios = tweet_freq
labels = target_classes
explode = [0.1, 0, 0]
colors = ['c','y','g']
# rotate so that first wedge is split by the x-axis
angle = -272 * overall_ratios[2]
wedges, *_ = ax1.pie(overall_ratios, autopct=lambda pct: func(pct,
    ↪overall_ratios),shadow=True,labels=labels,
    ↪explode=explode,colors=colors,startangle=angle)

# bar chart parameters
negative_sentiment_ratios = [airlines_sentiments_groups[x]['negative'] for x in
    ↪list_of_airlines]
bottom = 5
width = .2

# Adding from the top matches the legend.
for j, (height, label) in enumerate(reversed([*zip(negative_sentiment_ratios,
    ↪list_of_airlines)])):
    bottom -= height
    bc = ax2.bar(0, height, width, bottom=bottom, color='C0',
    ↪label=label,alpha=0.1 + 0.17 * j)
    ax2.bar_label(bc, labels=[height], label_type='center')

ax2.legend(loc=3)
ax2.set_title('Negative sentiments count per airline')
ax2.axis('off')
ax2.set_xlim(- 3.5 * width, 3.5 * width)

# use ConnectionPatch to draw lines between the two plots
theta1, theta2 = wedges[0].theta1, wedges[0].theta2
center, r = wedges[0].center, wedges[0].r
bar_height = sum(negative_sentiment_ratios)

# draw top connecting line
x = r * np.cos(np.pi / 180 * theta2) + center[0]
y = r * np.sin(np.pi / 180 * theta2) + center[1]
con = ConnectionPatch(xyA=(-width / 2, 0), coordsA=ax2.transData,

```

```

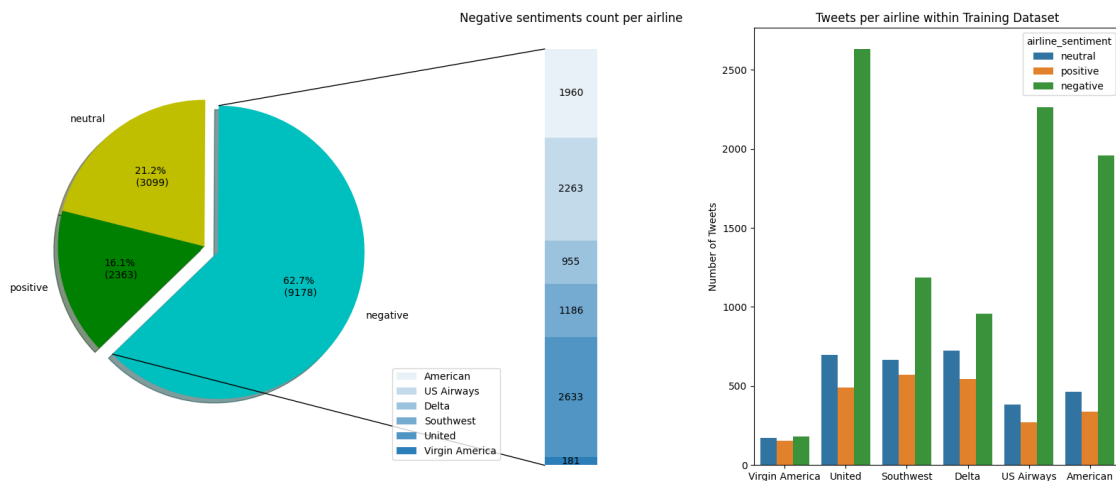
        xyB=(x, y), coordsB=ax1.transData)
con.set_color([0, 0, 0])
con.set_linewidth(1)
ax2.add_artist(con)

# draw bottom connecting line
x = r * np.cos(np.pi / 180 * theta1) + center[0]
y = r * np.sin(np.pi / 180 * theta1) + center[1]
con = ConnectionPatch(xyA=(-width / 2, -bar_height), coordsA=ax2.
    ↪transData,xyB=(x, y), coordsB=ax1.transData)
con.set_color([0, 0, 0])
ax2.add_artist(con)
con.set_linewidth(1)

ax3 = sns.countplot(x="airline", hue="airline_sentiment", data=dataset)
ax3.set_xlabel('')
ax3.set_ylabel('Number of Tweets')
ax3.set_title('Tweets per airline within Training Dataset')

plt.show()
del ax1,ax2,ax3

```



Date decomposition (feature creation).

```

[31]: def decode_date(date_str)-> pd.Series:
    aux = date_str.replace(' -0800', '')
    date_utc = datetime.strptime(aux, "%Y-%m-%d %H:%M:%S")
    return pd.Series([date_utc.weekday(), date_utc.day, date_utc.hour])

```

```
[32]: dataset[['week_day', 'day', 'hour']] = dataset.pop('tweet_created').
      ↪ apply(decode_date)
```

```
[33]: dataset.head(3)
```

```
[33]:      tweet_id  airline_sentiment  airline_sentiment_confidence \
0  570306133677760513             neutral                1.0000
1  570301130888122368             positive                0.3486
2  570301083672813571             neutral                0.6837


      negativereason  negativereason_confidence      airline \
0              N/A                NaN  Virgin America
1              N/A                0.0  Virgin America
2              N/A                NaN  Virgin America

      airline_sentiment_gold      name negativereason_gold  retweet_count \
0                NaN      cairdin                NaN            0
1                NaN      jnardino                NaN            0
2                NaN  yvonnalynn                NaN            0

      text tweet_coord \
0      @VirginAmerica What @dhepburn said.                NaN
1  @VirginAmerica plus you've added commercials t...                NaN
2  @VirginAmerica I didn't today... Must mean I n...                NaN

      tweet_location      user_timezone  week_day  day  hour
0                NaN  Eastern Time (US & Canada)      1   24   11
1                NaN  Pacific Time (US & Canada)      1   24   11
2      Lets Play  Central Time (US & Canada)      1   24   11
```


```
[34]: fig, axes = plt.subplots(figsize=(10,6))
      ax = sns.countplot(x="day", hue="airline_sentiment", data=dataset)
      ax.set_xlabel('Day')
      ax.set_ylabel('Number of Tweets')
      ax.set_title('Daily AirLine Tweets Distribution within a month')
      del ax,fig,axes
```



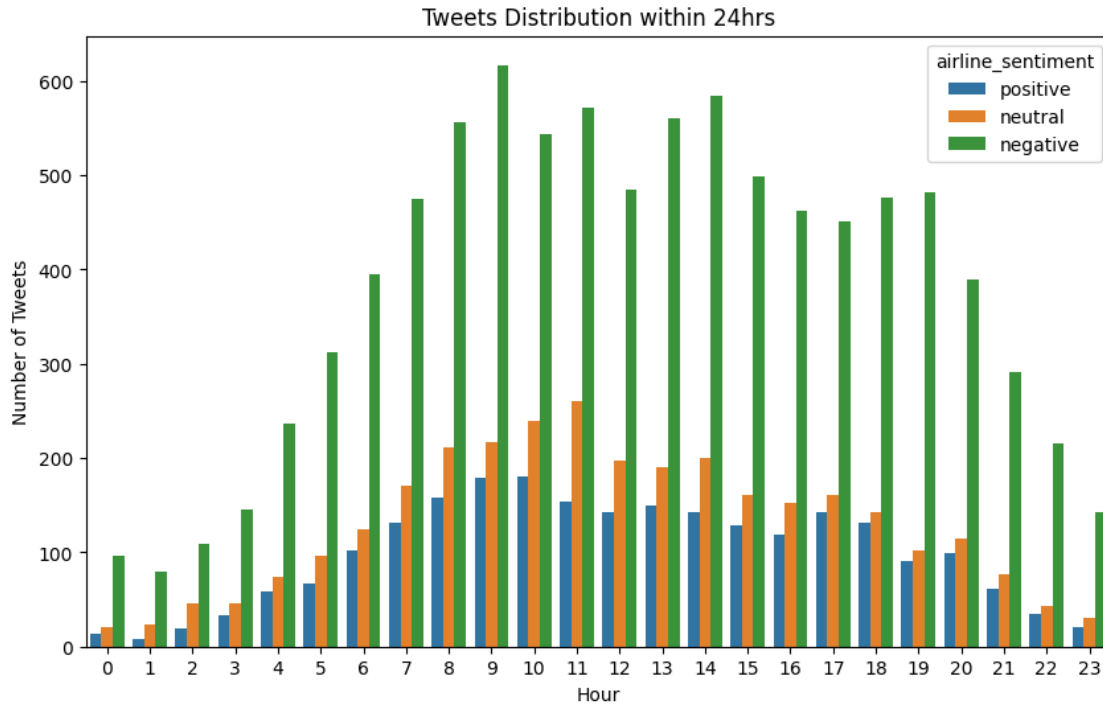
```
[35]: fig, axes = plt.subplots(figsize=(10,6))
ax = sns.countplot(x="week_day", hue="airline_sentiment", data=dataset)
ax.set_xlabel('Week day')
ax.
    ↪set(xticklabels=['Monday','Tuesday','Wednesday','Thursday','Friday','Saturday','Sunday'])
ax.set_ylabel('Number of Tweets')
ax.set_title('Daily AirLine Tweets Distribution within a week')
del ax,fig,axes
```

<ipython-input-35-64147deda210>:4: UserWarning: set_ticklabels() should only be used with a fixed number of ticks, i.e. after set_ticks() or using a FixedLocator.

```
ax.set(xticklabels=['Monday','Tuesday','Wednesday','Thursday','Friday','Saturday','Sunday'])
```



```
[36]: fig, axes = plt.subplots(figsize=(10,6))
      ax = sns.countplot(x="hour", hue="airline_sentiment", data=dataset)
      ax.set_xlabel('Hour')
      ax.set_ylabel('Number of Tweets')
      ax.set_title('Tweets Distribution within 24hrs')
      del ax,fig,axes
```



1.1.1 Top 10 tweet authors

```
[37]: top_tweeter = Counter([name for name in dataset['name']])
top_tweeter_df = pd.DataFrame(top_tweeter.most_common(10))
top_tweeter_df.columns = ['top_tweet_authors', 'count']
print("Top 10 Authors", '-'*15)
top_tweeter_df.style.background_gradient(cmap='inferno')
```

Top 10 Authors -----

```
[37]: <pandas.io.formats.style.Styler at 0x7824f30f3c10>
```

1.1.2 Top reasons for complaints.

```
[38]: top_complain = Counter([reason for reason in dataset[dataset['negativereason'] !=
    ↳ 'N/A']['negativereason']])
top_complain_df = pd.DataFrame(top_complain.most_common(10))
top_complain_df.columns = ['top_complain_reason', 'count']
print("Top 10 complain", '-'*15)
top_complain_df.style.background_gradient(cmap='inferno')
```

Top 10 complain -----

```
[38]: <pandas.io.formats.style.Styler at 0x7824f3144c70>
```


A closer look at the two top authors

```
[39]: dataset[dataset['name'] == 'JetBlueNews'][['airline', 'negativereason', 'airline_sentiment', 'retweet_count']].value_counts()
```

```
[39]: airline      negativereason  airline_sentiment  retweet_count
Delta      N/A                neutral              0              56
          Can't Tell          positive             0              5
          N/A                negative             0              1
Virgin America  N/A          neutral              0              1
Name: count, dtype: int64
```

```
[40]: dataset[dataset['name'] == 'kboasspotter'][['airline', 'negativereason', 'airline_sentiment', 'retweet_count']].value_counts()
```

```
[40]: airline      negativereason      airline_sentiment  retweet_count
Delta      N/A                neutral              0              22
          Canceled Flight      negative             0              6
          Customer Service Issue negative             0              1
          Flight Booking Problems negative             0              1
          Late Flight          negative             0              1
Name: count, dtype: int64
```

1.1.3 Data preprocessing, feature engineering.

```
[41]: #For removing user tags(@user)
def remove_user_tags(tweet:str)->tuple[str,int]:
    user = re.compile(r'@\S+')
    initial = len(tweet)
    text = user.sub('',tweet)
    result = ''.join([i for i in text if not i.isdigit()])
    final = len(result)
    number = 0
    for i in text:
        if i.isdigit():
            number+=1
    if( final == initial - number):
        return (result,0)
    else:
        return (result,1)

#For removing Url links
def remove_url(tweet:str)->tuple[str,int]:
    url = re.compile(r'https?://\S+|www\.\S+')
    initial = len(tweet)
```

```

text = url.sub(r'',tweet,re.IGNORECASE)
result = ''.join([i for i in text if not i.isdigit()])
final = len(result)
number = 0
for i in text:
    if i.isdigit():
        number += 1
if( final == initial - number):
    return (result,0)
else:
    return (result,1)

```

```

[42]: def clean_and_create_feature(tweet:str)-> pd.Series:
      (str_without_tags,tag_flag) = remove_user_tags(tweet)
      (str_without_links,url_flag) = remove_url(str_without_tags)
      return pd.Series([str_without_links,tag_flag,url_flag])

```

```

[43]: dataset[['text','user_tag','url_flag']] = dataset.apply(lambda x:
    ↪ clean_and_create_feature(x.text),axis=1)

```

```

[44]: dataset.head(3)

```

```

[44]:      tweet_id  airline_sentiment  airline_sentiment_confidence \
0  570306133677760513             neutral                1.0000
1  570301130888122368             positive                0.3486
2  570301083672813571             neutral                0.6837

      negativereason  negativereason_confidence      airline \
0              N/A                NaN  Virgin America
1              N/A                0.0  Virgin America
2              N/A                NaN  Virgin America

      airline_sentiment_gold      name negativereason_gold  retweet_count \
0                NaN      cairdin                NaN                0
1                NaN      jnardino                NaN                0
2                NaN      yvonnalynn                NaN                0

      text  tweet_coord \
0      What said.                NaN
1  plus you've added commercials to the experien...                NaN
2  I didn't today... Must mean I need to take an...                NaN

      tweet_location      user_timezone  week_day  day  hour  user_tag \
0                NaN  Eastern Time (US & Canada)      1  24  11      1
1                NaN  Pacific Time (US & Canada)      1  24  11      1
2      Lets Play  Central Time (US & Canada)      1  24  11      1

```

	url_flag
0	0
1	0
2	0

```
[45]: print(dataset.groupby("airline_sentiment", group_keys=True)[['user_tag']].
      ↪value_counts())
print('-'*35)
print(dataset.groupby("airline_sentiment", group_keys=True)[['url_flag']].
      ↪value_counts())
```

airline_sentiment	user_tag	
negative	1	9178
neutral	1	3099
positive	1	2363

Name: count, dtype: int64

```
-----
```

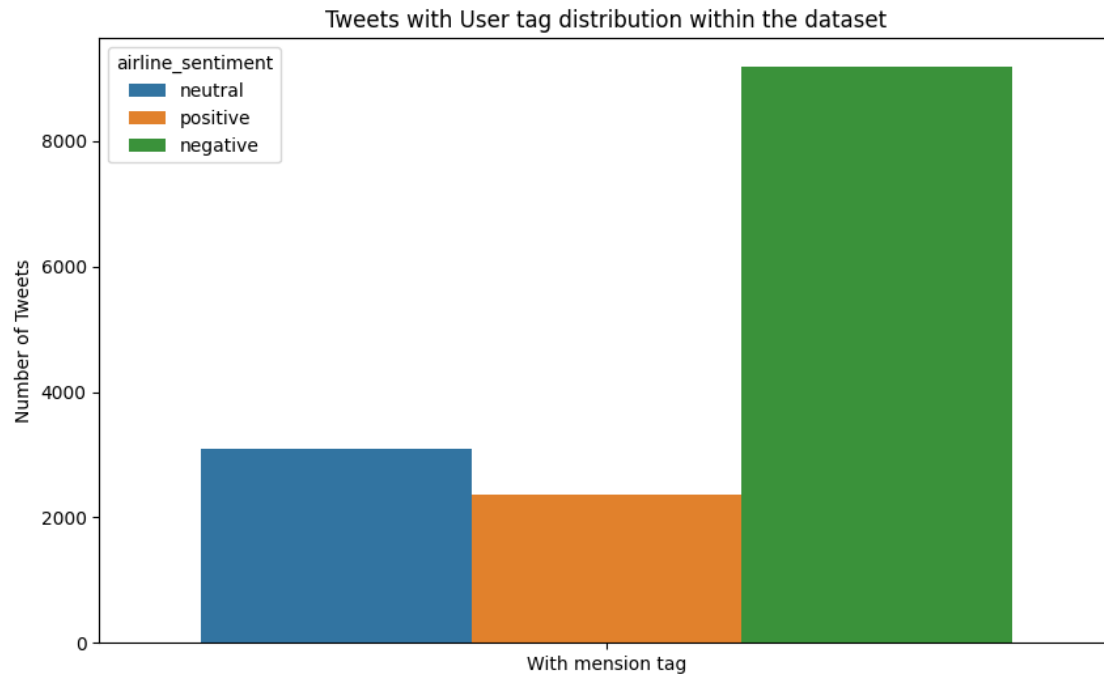
airline_sentiment	url_flag	
negative	0	8730
	1	448
neutral	0	2603
	1	496
positive	0	2134
	1	229

Name: count, dtype: int64

```
[46]: f, axes = plt.subplots(figsize=(10,6))
ax = sns.countplot(x="user_tag", hue="airline_sentiment", data=dataset)
ax.set(xticklabels=['With mension tag'])
ax.set_xlabel('')
ax.set_ylabel('Number of Tweets')
ax.set_title('Tweets with User tag distribution within the dataset')
del ax,f,axes
```

<ipython-input-46-05fe470843a3>:3: UserWarning: set_ticklabels() should only be used with a fixed number of ticks, i.e. after set_ticks() or using a FixedLocator.

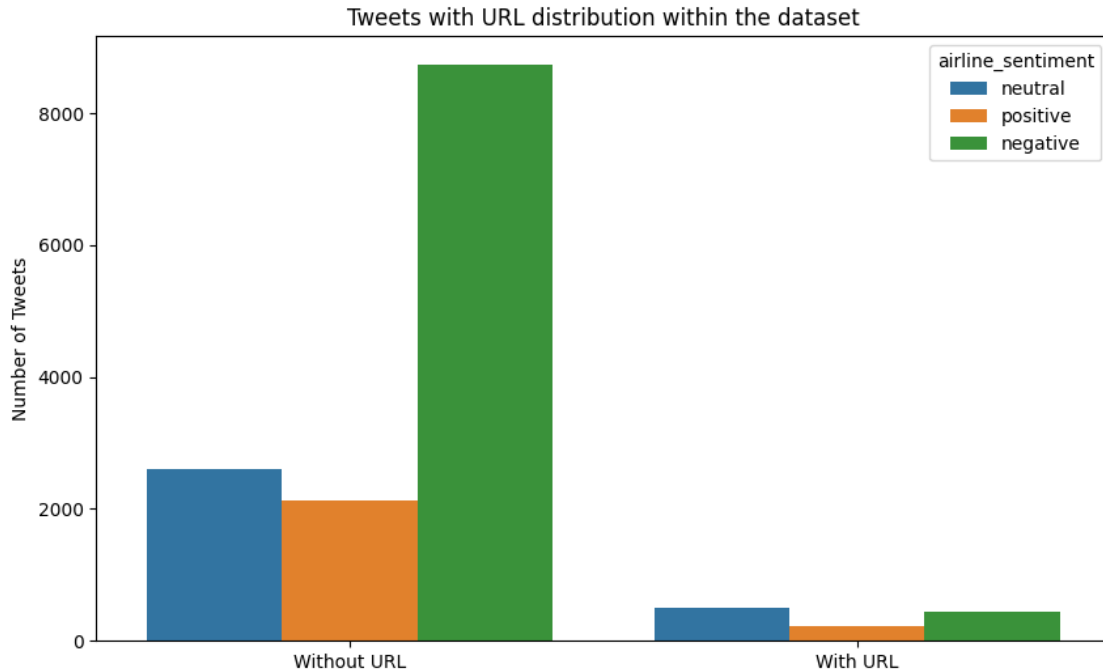
```
ax.set(xticklabels=['With mension tag'])
```



```
[47]: f, axes = plt.subplots(figsize=(10, 6))
      ax = sns.countplot(x="url_flag", hue="airline_sentiment", data=dataset)
      ax.set(xticklabels=['Without URL', 'With URL'])
      ax.set_xlabel('')
      ax.set_ylabel('Number of Tweets')
      ax.set_title('Tweets with URL distribution within the dataset')
      del ax, f, axes
```

<ipython-input-47-5d468f3940bb>:3: UserWarning: set_ticklabels() should only be used with a fixed number of ticks, i.e. after set_ticks() or using a FixedLocator.

```
ax.set(xticklabels=['Without URL', 'With URL'])
```



```
[48]: def happy_emoticons_removal(tweet:str) ->tuple[str,int]:
    happy = re.compile(r"([xX;:]-?[dDpP])")
    initial = len(tweet)
    text = happy.sub(r'',tweet)
    result = ''.join([i for i in text if not i.isdigit()])
    final = len(result)
    number = 0
    for i in text:
        if i.isdigit():
            number += 1
    if( final != initial - number):
        return (result,1)
    return (result,0)

def sad_emoticons_removal(tweet:str) ->tuple[str,int]:
    sad = re.compile(r"[:;](['\"]?[-~]?[/\(\|C<>{} \[\] +)")
    initial = len(tweet)
    text = sad.sub(r'',tweet)
    result = ''.join([i for i in text if not i.isdigit()])
    final = len(result)
    number = 0
    for i in text:
        if i.isdigit():
            number += 1
    if( final != initial - number):
```

```

    return (result,1)
    return (result,0)

```

```

[49]: def emoticon_removal_and_feature_creation(tweet:str)-> pd.Series:
      (str_without_happy_emoji,happy_emoji_flag) = happy_emoticons_removal(tweet)
      (str_without_emoji,sad_emoji_flag) =
      ↪sad_emoticons_removal(str_without_happy_emoji)
      return pd.Series([str_without_emoji,happy_emoji_flag,sad_emoji_flag])

```

```

[50]: dataset[['text','happy_emoji','sad_emoji']] = dataset.apply(lambda x:
      ↪emoticon_removal_and_feature_creation(x.text),axis=1)

```

```

[51]: f, axes = plt.subplots(1,2,figsize=(10, 6))
      ax = sns.countplot(x="happy_emoji", hue="airline_sentiment",
      ↪data=dataset,ax=axes[0])
      ax.set(xticklabels=['False','True'])
      ax.set_xlabel('"Happy" emoji on Tweet')
      ax.set_ylabel('Number of Tweets')
      ax = sns.countplot(x="happy_emoji", hue="airline_sentiment",
      ↪data=dataset[dataset.happy_emoji==1],ax=axes[1])
      ax.set(xticklabels=['True'])
      ax.set_xlabel('"Happy" emoji on Tweet')
      ax.set_ylabel('Number of Tweets')
      f.suptitle('Distribution of "Happy" emoji within the dataset',fontsize=15)
      del ax,f,axes

```


<ipython-input-51-91fc4191489b>:3: UserWarning: set_ticklabels() should only be used with a fixed number of ticks, i.e. after set_ticks() or using a FixedLocator.

```
ax.set(xticklabels=['False','True'])
```

<ipython-input-51-91fc4191489b>:7: UserWarning: set_ticklabels() should only be used with a fixed number of ticks, i.e. after set_ticks() or using a FixedLocator.

```
ax.set(xticklabels=['True'])
```

Distribution of "Happy" emoji within the dataset



```
[52]: f, axes = plt.subplots(1,2,figsize=(10, 6))
ax = sns.countplot(x="sad_emoji", hue="airline_sentiment",
    ↳data=dataset,ax=axes[0])
ax.set(xticklabels=['False','True'])
ax.set_xlabel('"Sad" emoji on Tweet')
ax.set_ylabel('Number of Tweets')
ax = sns.countplot(x="sad_emoji", hue="airline_sentiment", data=dataset[dataset.
    ↳sad_emoji==1],ax=axes[1])
ax.set(xticklabels=['True'])
ax.set_xlabel('"Sad" emoji on Tweet')
ax.set_ylabel('Number of Tweets')
f.suptitle('Distribution of "Sad" emoji within the dataset',fontsize=15)
del ax,f,axes
```


<ipython-input-52-52f20b644c5b>:3: UserWarning: set_ticklabels() should only be used with a fixed number of ticks, i.e. after set_ticks() or using a FixedLocator.

```
ax.set(xticklabels=['False','True'])
```

<ipython-input-52-52f20b644c5b>:7: UserWarning: set_ticklabels() should only be used with a fixed number of ticks, i.e. after set_ticks() or using a FixedLocator.

```
ax.set(xticklabels=['True'])
```

Distribution of "Sad" emoji within the dataset



```
[53]: print(dataset.groupby("airline_sentiment", group_keys=True)[['happy_emoji']].
      ↪value_counts())
print('-'*40)
print(dataset.groupby("airline_sentiment", group_keys=True)[['sad_emoji']].
      ↪value_counts())
```

```
airline_sentiment  happy_emoji
negative          0           8707
                  1           471
neutral           0          2962
                  1           137
positive          0          2128
                  1           235
Name: count, dtype: int64
```

```
-----
airline_sentiment  sad_emoji
negative          0          9073
                  1           105
neutral           0          3078
                  1            21
positive          0          2358
                  1             5
Name: count, dtype: int64
```



```
[54]: positive_sentiment = dataset[dataset.airline_sentiment == "positive"]
positive_text=positive_sentiment['text']
negative_sentiment = dataset[dataset.airline_sentiment == 'negative']
negative_text=negative_sentiment['text']
neutral_sentiment = dataset[dataset.airline_sentiment == 'neutral']
neutral_text=neutral_sentiment['text']
complain_text = top_complain_df['top_complain_reason']
top_authors = top_tweeter_df["top_tweet_authors"]

[55]: # Create and generate a word
fig, ax = plt.subplots(1, 5, figsize=(25, 8),edgecolor = 'k')
positive_tweet = WordCloud(width = 200,height = 300,
    colormap="Paired",background_color = 'black',max_words = 90,stopwords = STOPWORDS).generate(str(positive_text))
negative_tweet = WordCloud(width = 200,height = 300,
    colormap="Paired",background_color = 'black',max_words = 90,stopwords = STOPWORDS).generate(str(negative_text))
neutral_tweet = WordCloud(width = 200,height = 300,
    colormap="Paired",background_color = 'black',max_words = 90,stopwords = STOPWORDS).generate(str(neutral_text))
complain_tweet = WordCloud(width = 200,height = 300,
    colormap="Paired",background_color = 'black',max_words = 90,stopwords = STOPWORDS).generate(str(complain_text))
top_authors_tweet = WordCloud(width = 200,height = 300,
    colormap="Paired",background_color = 'black',max_words = 90,stopwords = STOPWORDS).generate(str(top_authors))

ax[0].imshow(positive_tweet)
ax[0].axis('off')
ax[0].set_title('Positive Sentiment')

ax[1].imshow(negative_tweet)
ax[1].axis('off')
ax[1].set_title('Negative Sentiment')

ax[2].imshow(neutral_tweet)
ax[2].axis('off')
ax[2].set_title('Neutral Sentiment')

ax[3].imshow(complain_tweet)
ax[3].axis('off')
ax[3].set_title('Top complain from clients')

ax[4].imshow(top_authors_tweet)
```


1.2 Drop less beneficial columns

```
[60]: dataset.  
      ↪drop(['user_tag', 'tweet_id', 'name', 'negativereason', 'negativereason_confidence', 'airline_se
```

```
[61]: print("Sum of missing values", sum(dataset.isnull().sum()))
```

Sum of missing values 0

1.2.1 y_i can take values between 0 to $N - 1$ categorories

```
[62]: target_dict = {'positive':1, 'negative': 0, 'neutral': 2}  
      print(target_dict)  
      dataset['target'] = dataset['airline_sentiment'].map(target_dict)
```

{'positive': 1, 'negative': 0, 'neutral': 2}

```
[63]: dataset.head(3)
```

```
[63]:  airline_sentiment  airline_sentiment_confidence  retweet_count  \  
0          neutral                1.0000                0  
1        positive                0.3486                0  
2          neutral                0.6837                0  
  
      text  week_day  day  hour  url_flag  \  
0      said         1   24   11         0  
1  plu ad commerci eerien  tacki         1   24   11         0  
2 today must mean need take anoth trip         1   24   11         0  
  
      happy_emoji  sad_emoji  airline_code  target  
0              0          0              0        2  
1              1          0              0        1  
2              0          0              0        2
```

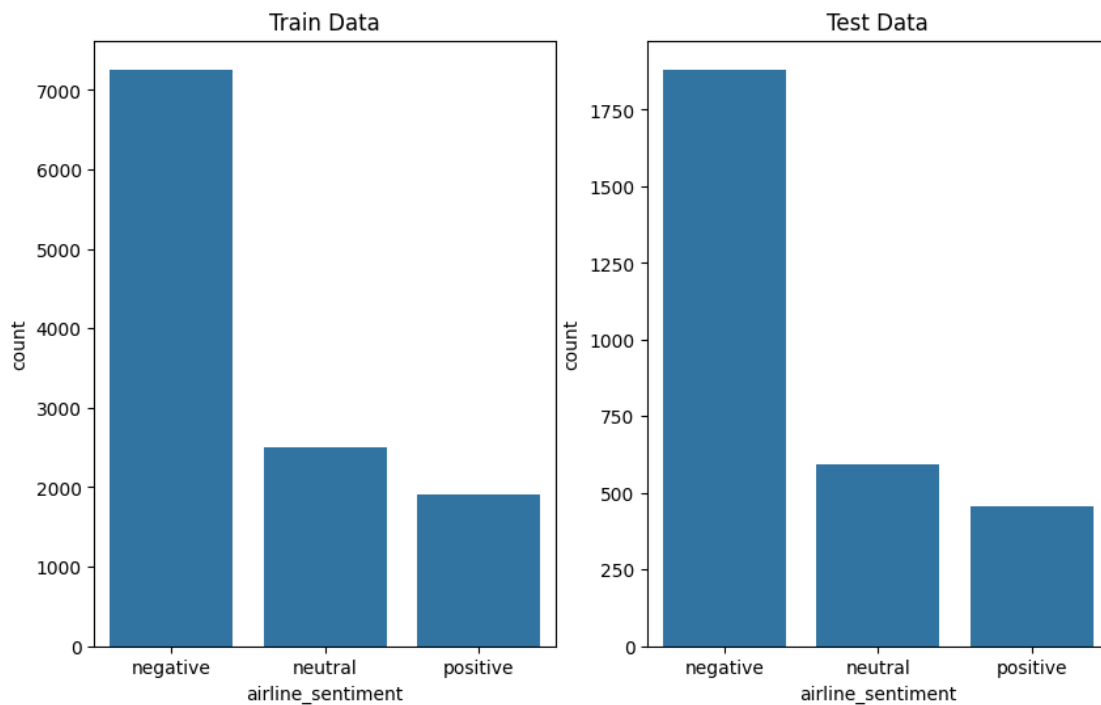
```
[64]: base_line_df = dataset.copy()
```

```
[65]: def split_dataset(df:pd.DataFrame, test_percentage:float)-> tuple[pd.  
      ↪DataFrame, pd.DataFrame]:  
      shuffle = np.random.permutation(len(df))  
      test_size = int(len(df) * test_percentage)  
      test_aux = shuffle[:test_size]  
      train_aux = shuffle[test_size:]  
      return (df.iloc[train_aux], df.iloc[test_aux])
```

```
[66]: train, test = split_dataset(base_line_df, 0.2)
```

```
[67]: train.drop_duplicates(inplace=True)
```

```
[68]: fig, ax = plt.subplots(1, 2, figsize=(10, 6))
sns.countplot(x='airline_sentiment', data=train, ax=ax[0])
sns.countplot(x='airline_sentiment', data=test, ax=ax[1])
ax[0].set_title('Train Data')
ax[1].set_title('Test Data')
plt.show()
del ax,fig
```



```
[69]: print(train.shape, test.shape)
```

```
(11659, 12) (2928, 12)
```

```
[70]: print(train[['target', 'airline_sentiment']].value_counts())
print("test samples", '-'*20)
print(test[['target', 'airline_sentiment']].value_counts())
```

```
target  airline_sentiment
0      negative           7260
2      neutral           2498
1      positive           1901
Name: count, dtype: int64
test samples -----
target  airline_sentiment
0      negative           1882
2      neutral            592
```

```
1         positive          454
Name: count, dtype: int64
```

```
[71]: train.head(3)
```

```
[71]:      airline_sentiment  airline_sentiment_confidence  retweet_count  \
4219          negative                1.0000                0
13090         negative                0.6326                0
9033          negative                1.0000                0

      text  week_day  day  hour  \
4219  flight already cancel flightl tri get home tw...      1   17   11
13090  realli appreci great custom servic one servic...      0   23   11
9033   spoke someon told breast pump medic equip pla...      1   24    9

      url_flag  happy_emoji  sad_emoji  airline_code  target
4219         0           0           0             1        0
13090         0           0           0             5        0
9033         0           0           0             4        0
```

1.2.2 Baseline Model Implementation

```
[72]: y = train['target'].to_numpy()
```

```
[73]: y_t = test['target'].to_numpy()
```

1.2.3 Bag of words (BoW).

```
[74]: vec = TfidfVectorizer(ngram_range=(1,2),stop_words='english').fit(train['text'])
```

```
[75]: X = vec.transform(train['text'])
```

```
[76]: X_t =vec.transform(test['text'])
```

```
[77]: vocab_size = len(vec.vocabulary_) + 1
print("Vocabulary Size :", vocab_size)
```

Vocabulary Size : 61691

Linear support vector machine (With class weight to compensate for the imbalance).

```
[78]: param_grid = {'C': [0.01, 0.1, 1.0, 10.0, 100.0]}
clf = LinearSVC(loss='hinge',class_weight="balanced")
```

```
[79]: grids = GridSearchCV(clf, param_grid,verbose=1,n_jobs=-1)
grids = grids.fit(X, y)
print ("Best parameters: %s" % grids.best_params_)
```

Fitting 5 folds for each of 5 candidates, totalling 25 fits
Best parameters: {'C': 1.0}

```
[80]: base_line_predictions = grids.predict(X_t)
```

```
[81]: def score_board_df(y_test, predictions, model_name, averaging_method='macro')-> pd.
      ↪ DataFrame:
      f1 = "%0.3f" % f1_score(y_test, predictions, average=averaging_method)
      prec = "%0.3f" % precision_score(y_test, predictions, average=averaging_method)
      rec = "%0.3f" % recall_score(y_test, predictions, average=averaging_method)
      acc = "%0.3f" % accuracy_score(y_test, predictions)
      return pd.DataFrame({"classifier-name": model_name, "f1_score": [f1], "precision":
      ↪ [prec], "recall": [rec], "accuracy": [acc], "Average-method": averaging_method})
```


```
[262]: score_board = score_board_df(y_t, base_line_predictions, "LinearSVC +
      ↪ Weights", "weighted")
```

```
[263]: score_board
```

```
[263]:      classifier-name  f1_score  precision  recall  accuracy  Average-method
0  LinearSVC + Weights    0.785    0.783  0.789    0.789    weighted
```

```
[264]: def conf_matrix(y_true, y_pred)-> pd.DataFrame:
      # Creating a confusion matrix
      cm = confusion_matrix(y_true, y_pred)
      con_mat = pd.DataFrame(cm, index=np.unique(y_true), columns=np.
      ↪ unique(y_true))
      #Plotting the confusion matrix
      plt.figure(figsize=(10,6))
      ax = sns.heatmap(con_mat, annot=True, annot_kws={"size": 16}, fmt='g',
      ↪ cmap=plt.cm.Blues, cbar=False)
      ax.set_title('Confusion matrix')
      ax.set_xlabel('\nPredicted labels')
      ax.set_ylabel('True labels')
      plt.show()
      del ax
      return con_mat
```

```
[265]: _ = conf_matrix(y_t, base_line_predictions)
```



```
[266]: #print(test[y_t!=base_line_predictions]) #misclassified texts
```

1.2.4 Word Embedding

1.2.5 Model 1 (Word2vec Embedding)

```
[267]: def create_word_vector(text):
        tknzs = TweetTokenizer()
        tweet = tknzs.tokenize(text)
        aux = [int(vec.vocabulary_[k]) for k in tweet if k in vec.vocabulary_.keys()]
        return aux
```

```
[268]: model_1_train = train.copy()
        model_1_test = test.copy()
```

```
[269]: model_1_train['text'] = model_1_train.text.apply(create_word_vector)
        model_1_test['text'] = model_1_test.text.apply(create_word_vector)
```

```
[270]: MAX_LENGTH = len(max(model_1_train.text, key=len))
        print(MAX_LENGTH)
```

```
[271]: # plot the distribution of review lengths
#sns.distplot([len(x) for x in model_1_train.text])
#plt.xlim([0, 256]);
#plt.xlabel('Token count')
```

```
[272]: X_train= pad_sequences(model_1_train.text, maxlen=MAX_LENGTH,padding='post')
```

```
[273]: X_test = pad_sequences(model_1_test.text, maxlen=MAX_LENGTH,padding='post')
```

```
[274]: Y_train = keras.utils.to_categorical(y, num_classes=3)
```

```
[275]: Y_test = keras.utils.to_categorical(y_t, num_classes=3)
```

```
[276]: Y_train[27:30]
```

```
[276]: array([[1., 0., 0.],
         [0., 0., 1.],
         [0., 1., 0.]])
```

```
[277]: model_1_train.airline_sentiment[27:30]
```

```
[277]: 14026    negative
      9324     neutral
      11389    positive
      Name: airline_sentiment, dtype: object
```

1.2.6 Model 1 RNN

```
[278]: METRICS = [keras.metrics.CategoricalAccuracy(name='accuracy'),keras.metrics.
    ↳Precision(name='precision'),keras.metrics.Recall(name='recall'),keras.
    ↳metrics.AUC(name='prc', curve='PR')]
embedding_vector_length= 32
```

```
[279]: def create_model(vocab_size:int,embedding_vector_length:int,weights:
    ↳list=None,metrics:list=METRICS,trainable:bool=False):
    model = keras.models.Sequential()
    model.add(keras.layers.
    ↳Embedding(vocab_size,embedding_vector_length,weights=weights,trainable=trainable))
    model.add(keras.layers.Bidirectional(keras.layers.
    ↳LSTM(64,return_sequences=True,dropout=0.2, recurrent_dropout=0.2)))
    model.add(keras.layers.GlobalAveragePooling1D())
    model.add(keras.layers.Dense(200, activation='relu'))
    model.add(keras.layers.Dense(3, activation='softmax'))
    model.compile(loss=keras.losses.CategoricalCrossentropy(),optimizer=tf.keras.
    ↳optimizers.Adam(learning_rate=0.005),metrics=METRICS)
    return model
```



```
[280]: model1 = create_model(vocab_size,embedding_vector_length)
print(model1.summary())
```

Model: "sequential_11"

Layer (type) ↳Param #	Output Shape	
embedding_13 (Embedding) ↳(unbuilt)	?	0
bidirectional_11 (Bidirectional) ↳(unbuilt)	?	0
global_average_pooling1d_11 ↳(unbuilt) (GlobalAveragePooling1D) ↳	?	0
dense_26 (Dense) ↳(unbuilt)	?	0
dense_27 (Dense) ↳(unbuilt)	?	0

Total params: 0 (0.00 B)

Trainable params: 0 (0.00 B)

Non-trainable params: 0 (0.00 B)

None

```
[281]: EPOCHS = 100
BATCH_SIZE = 2048
```

```
[282]: steps_per_epoch = int(np.ceil(3.0*tweet_freq[0]/BATCH_SIZE))
```

```
[283]: lr = tf.keras.callbacks.ReduceLROnPlateau(factor = 0.1, min_lr = 0.01,
↳monitor='val_prc')
early_stopping = tf.keras.callbacks.
↳EarlyStopping(monitor='val_prc',mode='max',patience=10,restore_best_weights=True)
```

```

model_1_history = model1.
    ↪fit(X_train,Y_train,batch_size=BATCH_SIZE,epochs=EPOCHS,callbacks=[lr,early_stopping],valid
    ↪,Y_test),verbose=0)

```

```

[284]: def plot_metrics(history):
    metrics = ['loss', 'prc', 'precision', 'recall']
    for n, metric in enumerate(metrics):
        name = metric.replace("_", " ").capitalize()
        plt.subplot(2,2,n+1)
        plt.rcParams["figure.figsize"] = (10,6)
        plt.plot(history.epoch, history.history[metric], label='Train')
        plt.plot(history.epoch, history.history['val_'+metric], linestyle="--",
    ↪label='Val')
        plt.xlabel('Epoch')
        plt.ylabel(name)
        if metric == 'loss':
            plt.ylim([0, plt.ylim()[1]])
        elif metric == 'auc':
            plt.ylim([0.8,1])
        else:
            plt.ylim([0,1])

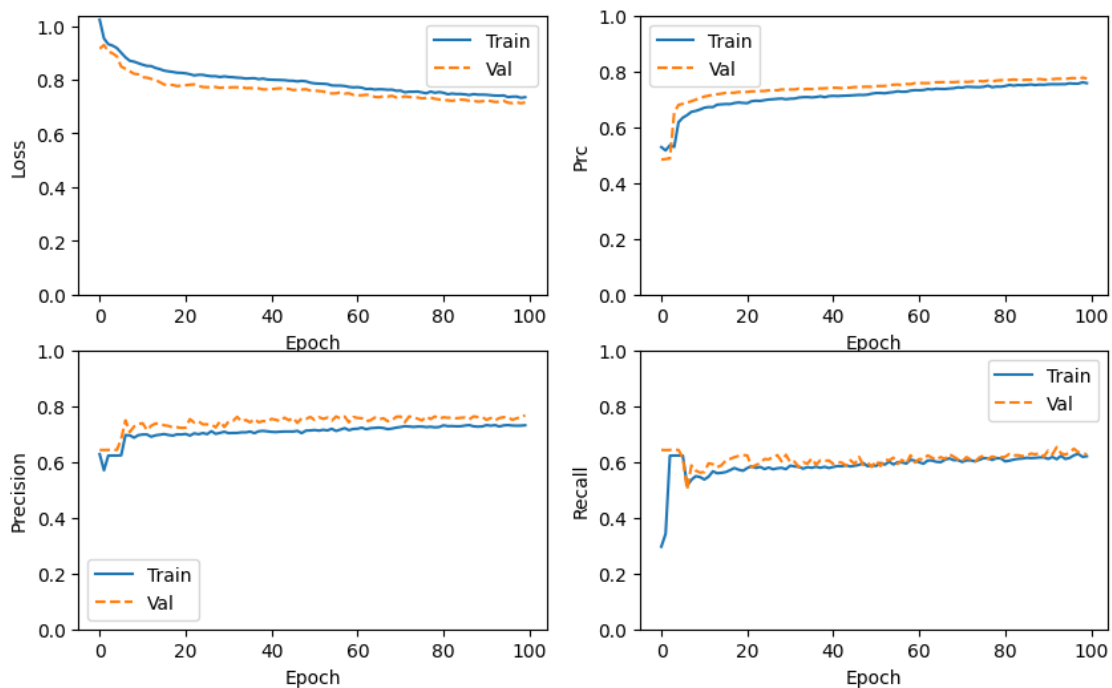
        plt.legend()

```

```

[285]: _ = plot_metrics(model_1_history)

```



```
[286]: model_1_results = model1.evaluate(X_test ,Y_test,verbose=0)
```


```
[287]: for name, value in zip(model1.metrics_names, model_1_results):
        print(name, ': {:.3f}'.format(value))
        print()
```

loss : 0.713

compile_metrics : 0.713

```
[288]: model_1_predictions = np.argmax(model1.predict(X_test,verbose=0), axis=-1)
```

```
[289]: _ = conf_matrix(y_t, model_1_predictions)
```



```
[290]: score_board2 = score_board_df(y_t, model_1_predictions,"RNN-Word2Vec")
```

```
[291]: score_board = pd.concat([score_board,score_board2],ignore_index=True,
        ↪sort=False)
        score_board
```

```
[291]: classifier-name f1_score precision recall accuracy Average-method
0 LinearSVC + Weights 0.785 0.783 0.789 0.789 weighted
```

1	RNN-Word2Vec	0.585	0.647	0.557	0.713	macro
---	--------------	-------	-------	-------	-------	-------

1.2.7 Model 2 (With class Weights)

Calculate class weights

```
[292]: weight_for_0 = (1/tweet_freq[0]) * (sum(tweet_freq) / 3.0)
weight_for_1 = (1/tweet_freq[2]) * (sum(tweet_freq) / 3.0)
weight_for_2 = (1/tweet_freq[1]) * (sum(tweet_freq) / 3.0)
class_weights = {0: weight_for_0, 1: weight_for_1, 2: weight_for_2}
print("Weight for class negative: {:.2f}".format(weight_for_0))
print("Weight for class positive: {:.2f}".format(weight_for_1))
print("Weight for class neutral: {:.2f}".format(weight_for_2))
```

Weight for class negative: 0.53

Weight for class positive: 2.07

Weight for class neutral: 1.57

```
[293]: model2 = create_model(vocab_size,32)
print(model2.summary())
```

Model: "sequential_12"

Layer (type)	Output Shape	
Param #		
embedding_14 (Embedding)	?	0
(unbuilt)		
bidirectional_12 (Bidirectional)	?	0
(unbuilt)		
global_average_pooling1d_12	?	0
(unbuilt)		
(GlobalAveragePooling1D)		
dense_28 (Dense)	?	0
(unbuilt)		
dense_29 (Dense)	?	0
(unbuilt)		

Total params: 0 (0.00 B)

Trainable params: 0 (0.00 B)


Non-trainable params: 0 (0.00 B)

None

```
[294]: lr = tf.keras.callbacks.ReduceLROnPlateau(factor = 0.1, min_lr = 0.01,  
        ↪monitor='val_loss')  
early_stopping = tf.keras.callbacks.  
        ↪EarlyStopping(monitor='val_prc',patience=8,restore_best_weights=True)  
model_2_history = model2.fit(X_train,Y_train,batch_size=int(BATCH_SIZE*0.  
        ↪5),epochs=int(EPOCHS*0.  
        ↪5),callbacks=[lr,early_stopping],validation_data=(X_test,  
        ↪Y_test),verbose=0,class_weight=class_weights)
```


```
[295]: model_2_results = model2.evaluate(X_test ,Y_test,verbose=0)
```

```
[296]: _ = plot_metrics(model_2_history)
```



```
[297]: for name, value in zip(model2.metrics_names, model_2_results):  
        print(name, ': {:.3f}'.format(value))  
print()  
model_2_predictions = np.argmax(model2.predict(X_test,verbose=0), axis=1)  
_ = conf_matrix(y_t, model_2_predictions)
```

```
loss : 0.912
compile_metrics : 0.643
```



```
[298]: score_board3 = score_board_df(y_t, model_2_predictions, "RNN-Word2Vec +
↪Weights", "weighted")
score_board = pd.concat([score_board, score_board3], ignore_index=True,
↪sort=False)
score_board
```

```
[298]: classifier-name f1_score precision recall accuracy Average-method
0 LinearSVC + Weights 0.785 0.783 0.789 0.789 weighted
1 RNN-Word2Vec 0.585 0.647 0.557 0.713 macro
2 RNN-Word2Vec + Weights 0.649 0.657 0.643 0.643 weighted
```

1.2.8 Model 3 (Pre-Trained GloVe Embedding)

```
[299]: !wget http://nlp.stanford.edu/data/glove.6B.zip
!unzip glove.6B.zip
```

```
--2024-11-25 13:26:25-- http://nlp.stanford.edu/data/glove.6B.zip
Resolving nlp.stanford.edu (nlp.stanford.edu)... 171.64.67.140
Connecting to nlp.stanford.edu (nlp.stanford.edu)|171.64.67.140|:80...
connected.
```

```

HTTP request sent, awaiting response... 302 Found
Location: https://nlp.stanford.edu/data/glove.6B.zip [following]
--2024-11-25 13:26:25-- https://nlp.stanford.edu/data/glove.6B.zip
Connecting to nlp.stanford.edu (nlp.stanford.edu)|171.64.67.140|:443...
connected.
HTTP request sent, awaiting response... 301 Moved Permanently
Location: https://downloads.cs.stanford.edu/nlp/data/glove.6B.zip [following]
--2024-11-25 13:26:25-- https://downloads.cs.stanford.edu/nlp/data/glove.6B.zip
Resolving downloads.cs.stanford.edu (downloads.cs.stanford.edu)... 171.64.64.22
Connecting to downloads.cs.stanford.edu
(downloads.cs.stanford.edu)|171.64.64.22|:443... connected.
HTTP request sent, awaiting response... 200 OK
Length: 862182613 (822M) [application/zip]
Saving to: 'glove.6B.zip.3'

```

```

glove.6B.zip.3      100%[=====>] 822.24M  5.01MB/s    in 2m 39s

```

```

2024-11-25 13:29:04 (5.17 MB/s) - 'glove.6B.zip.3' saved [862182613/862182613]

```

```

Archive: glove.6B.zip
replace glove.6B.50d.txt? [y]es, [n]o, [A]ll, [N]one, [r]ename:

```

```
[300]: embeddings_index = dict()
```

```
[301]: f = open('glove.6B.100d.txt')
```

```
[302]: EMBEDDING_DIM = 100
```

```
[303]: for line in f:
        values = line.split()
        word = values[0]
        coefs = np.asarray(values[1:], dtype='float32')
        embeddings_index[word] = coefs
    f.close()
    print('Loaded %s word vectors.' % len(embeddings_index))

```

```

Loaded 400000 word vectors.

```

```
[304]: embedding_matrix = np.zeros((vocab_size, EMBEDDING_DIM))
        num_words_in_embedding = 0
        for word, i in vec.vocabulary.items():
            embedding_vector = embeddings_index.get(word)
            if embedding_vector is not None:
                num_words_in_embedding += 1
                embedding_matrix[i] = embedding_vector

```

```
[305]: model3 = create_model(embedding_matrix.shape[0], embedding_matrix.
    ↪shape[1], weights = [embedding_matrix])

```

```
print(model2.summary())
```

Model: "sequential_12"

Layer (type) ↳Param #	Output Shape	
embedding_14 (Embedding) ↳1,974,112	(None, 20, 32)	↳
bidirectional_12 (Bidirectional) ↳49,664	(None, 20, 128)	↳
global_average_pooling1d_12 ↳ 0 (GlobalAveragePooling1D) ↳	(None, 128)	↳
dense_28 (Dense) ↳25,800	(None, 200)	↳
dense_29 (Dense) ↳603	(None, 3)	↳

Total params: 2,202,315 (8.40 MB)

Trainable params: 76,067 (297.14 KB)

Non-trainable params: 1,974,112 (7.53 MB)


Optimizer params: 152,136 (594.29 KB)

None

```
[306]: lr = tf.keras.callbacks.ReduceLROnPlateau(factor = 0.1, min_lr = 0.01,↳  
↳monitor='val_prc')  
early_stopping = tf.keras.callbacks.  
↳EarlyStopping(monitor='val_prc',mode='max',patience=10,restore_best_weights=True)  
model_3_history = model3.fit(X_train,Y_train,batch_size=int(BATCH_SIZE*0.  
↳5),epochs=int(EPOCHS*0.  
↳5),callbacks=[lr,early_stopping],validation_data=(X_test,↳  
↳Y_test),verbose=0,class_weight=class_weights)
```



```
[307]: _ = plot_metrics(model_3_history)
```




```
[308]: model_3_results = model3.evaluate(X_test ,Y_test,verbose=0)
```

```
[309]: for name, value in zip(model3.metrics_names, model_2_results):
        print(name, ': {:.3f}'.format(value))
        print()
```

```
loss : 0.912
compile_metrics : 0.643
```

```
[310]: model_3_predictions = np.argmax(model3.predict(X_test,verbose=0), axis=-1)
```

```
[311]: _ = conf_matrix(y_t, model_3_predictions)
```



```
[312]: score_board4 = score_board_df(y_t, model_3_predictions, "RNN-Glove +  

↳Weights", "weighted")
```

```
[313]: score_board = pd.concat([score_board, score_board4], ignore_index=True,  

↳sort=False)  

score_board
```

```
[313]:
```

	classifier-name	f1_score	precision	recall	accuracy	Average-method
0	LinearSVC + Weights	0.785	0.783	0.789	0.789	weighted
1	RNN-Word2Vec	0.585	0.647	0.557	0.713	macro
2	RNN-Word2Vec + Weights	0.649	0.657	0.643	0.643	weighted
3	RNN-Glove + Weights	0.735	0.764	0.723	0.723	weighted

1.2.9 Model 4 (With multiple features - x_n)

```
[314]: other_features_train = model_1_train.  

↳drop(['target', 'text', 'airline_sentiment'], inplace=False, axis=1)
```

```
[315]: other_features_test = model_1_test.  

↳drop(['target', 'text', 'airline_sentiment'], inplace=False, axis=1)
```

```
[316]: other_features_train.head(3)
```

```
[316]:      airline_sentiment_confidence  retweet_count  week_day  day  hour  \
4219                        1.0000                0         1   17   11
13090                       0.6326                0         0   23   11
9033                        1.0000                0         1   24    9

      url_flag  happy_emoji  sad_emoji  airline_code
4219         0           0          0             1
13090         0           0          0             5
9033         0           0          0             4
```

```
[317]: other_features_train.shape
```

```
[317]: (11659, 9)
```

```
[318]: X_train_2 = np.asarray(other_features_train.values, dtype=int)
```

```
[319]: X_test_2 = np.asarray(other_features_test.values, dtype=int)
```

```
[320]: print(X_train.shape, X_train_2.shape)
```

```
(11659, 20) (11659, 9)
```

```
[321]: input_dense = keras.layers.Input(shape=(9,))
input_embedding = keras.layers.Input(shape=(MAX_LENGTH,))
embedding = keras.layers.Embedding(embedding_matrix.shape[0], embedding_matrix.
    ↪shape[1], weights=[embedding_matrix])(input_embedding)
lstm = keras.layers.LSTM(200)(embedding)
concat = keras.layers.Concatenate()([lstm, input_dense])
Dense1 = keras.layers.Dense(200, activation='relu')(concat)
Dropout = keras.layers.Dropout(0.2)(Dense1)
Output = keras.layers.Dense(3, activation='softmax')(Dropout)
model4 = keras.Model(inputs=[input_embedding, input_dense], outputs=[Output])
model4.compile(loss=tf.keras.losses.CategoricalCrossentropy(), optimizer=keras.
    ↪optimizers.Adam(learning_rate=0.005), metrics = METRICS)
print(model3.summary())
```

```
Model: "sequential_13"
```

Layer (type)	Output Shape	
↪Param #		
embedding_15 (Embedding)	(None, 20, 100)	↪
↪6,169,100		
bidirectional_13 (Bidirectional)	(None, 20, 128)	↪
↪84,480		

```

global_average_pooling1d_13      (None, 128)
↳ 0
(GlobalAveragePooling1D)
↳

dense_30 (Dense)                  (None, 200)
↳25,800

dense_31 (Dense)                  (None, 3)
↳603

```

Total params: 6,501,751 (24.80 MB)

Trainable params: 110,883 (433.14 KB)

Non-trainable params: 6,169,100 (23.53 MB)

Optimizer params: 221,768 (866.29 KB)

None

```

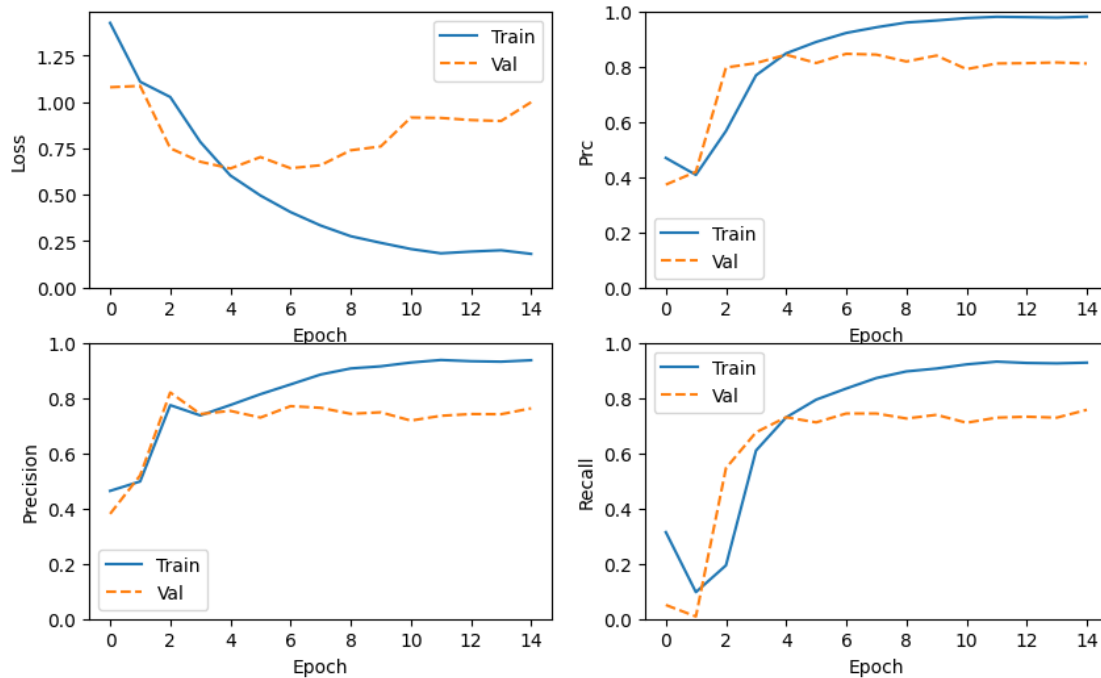
[322]: lr = tf.keras.callbacks.ReduceLROnPlateau(monitor='val_loss',
↳patience=10,factor=0.5, min_lr=0.001)
early_stop = keras.callbacks.
↳EarlyStopping(monitor='val_prc',mode='max',patience=8,
↳restore_best_weights=True)
model_4_history = model4.
↳fit([X_train,X_train_2],Y_train,batch_size=int(BATCH_SIZE*0.
↳25),epochs=EPOCHS,validation_data=([X_test, X_test_2], Y_test),
↳callbacks=[lr,early_stop],verbose=0,class_weight=class_weights)

```

```

[323]: plot_metrics(model_4_history)

```




```
[324]: model_4_results = model4.evaluate([X_test, X_test_2],Y_test,verbose=0)
```

```
[325]: for name, value in zip(model4.metrics_names, model_4_results):
        print(name, ': {:.3f}'.format(value))
        print()
```

```
loss : 0.643
compile_metrics : 0.761
```

```
[326]: model_4_predictions = np.argmax(model4.predict([X_test, X_test_2],verbose=0),
        ↪axis=-1)
```

```
[327]: _ = conf_matrix(y_t, model_4_predictions)
```



```
[328]: score_board5 = score_board_df(y_t, model_4_predictions, "RNN-Glove + Weights + \
      ↳ other features", "weighted")
```

```
[329]: score_board = pd.concat([score_board, score_board5], ignore_index=True, \
      ↳ sort=False)
score_board
```

```
[329]:
```

	classifier-name	f1_score	precision	recall	accuracy	\
0	LinearSVC + Weights	0.785	0.783	0.789	0.789	
1	RNN-Word2Vec	0.585	0.647	0.557	0.713	
2	RNN-Word2Vec + Weights	0.649	0.657	0.643	0.643	
3	RNN-Glove + Weights	0.735	0.764	0.723	0.723	
4	RNN-Glove + Weights + other features	0.769	0.787	0.761	0.761	

```
Average-method
0      weighted
1      macro
2      weighted
3      weighted
4      weighted
```

1.2.10 Model 5 (Experiment-Upsampling of minority classes using numpy)

```
[330]: train_labels = y.reshape(-1, 1)
       test_labels = y_t.reshape(-1, 1)

[331]: train_labels[27:30]

[331]: array([[0],
              [2],
              [1]])

[332]: train_features = X_train
       test_features = X_test

[333]: bool_train_labels_pos = ((train_labels[:, 0] != 0) & (train_labels[:, 0] != 2))

[334]: bool_train_labels_neg = ((train_labels[:, 0] != 1) & (train_labels[:, 0] != 2))

[335]: bool_train_labels_neu = ((train_labels[:, 0] != 0) & (train_labels[:, 0] != 1))

[336]: #bool_train_labels_neu[27:30]

[337]: pos_features = train_features[bool_train_labels_pos]
       pos_labels = train_labels[bool_train_labels_pos]

[338]: neg_features = train_features[bool_train_labels_neg]

[339]: neg_labels = train_labels[bool_train_labels_neg]

[340]: pos_ids = np.arange(len(pos_features))

[341]: random_choices1 = np.random.choice(pos_ids, len(neg_features))

[342]: res_pos_features = pos_features[random_choices1]
       res_pos_labels = pos_labels[random_choices1]

[343]: res_pos_features.shape

[343]: (7260, 20)

[344]: neu_features = train_features[bool_train_labels_neu]
       neu_labels = train_labels[bool_train_labels_neu]

[345]: neu_ids = np.arange(len(neu_features))

[346]: random_choices2 = np.random.choice(neu_ids, len(neg_features))
```

```
[347]: res_neu_features = neu_features[random_choices2]
       res_neu_labels = neu_labels[random_choices2]

[348]: res_neu_features.shape

[348]: (7260, 20)

[349]: resampled_features = np.concatenate([res_pos_features,res_neu_features,
      ↪neg_features], axis=0)

[350]: resampled_labels = np.concatenate([res_pos_labels,res_neu_labels, neg_labels],
      ↪axis=0)

[351]: order = np.arange(len(resampled_labels))
       np.random.shuffle(order)
       resampled_features = resampled_features[order]
       resampled_labels = resampled_labels[order]

[352]: resampled_features.shape

[352]: (21780, 20)

[353]: resampled_labels = keras.utils.to_categorical(resampled_labels, num_classes=3)

[354]: model5 = create_model(embedding_matrix.shape[0], embedding_matrix.
      ↪shape[1],weights = [embedding_matrix])
       print(model5.summary())
```

Model: "sequential_14"

Layer (type)	Output Shape	
↪Param #		
embedding_17 (Embedding)	?	
↪6,169,100		
bidirectional_14 (Bidirectional)	?	
↪(unbuilt)		0
global_average_pooling1d_14	?	0
↪(unbuilt)		
(GlobalAveragePooling1D)		
↪		
dense_34 (Dense)	?	0
↪(unbuilt)		

dense_35 (Dense)
↪(unbuilt)

?

0

Total params: 6,169,100 (23.53 MB)


Trainable params: 0 (0.00 B)

Non-trainable params: 6,169,100 (23.53 MB)

None

```
[355]: lr = tf.keras.callbacks.ReduceLROnPlateau(factor = 0.1, min_lr = 0.01, ↪  
↪monitor='val_prc')  
early_stopping = tf.keras.callbacks.  
↪EarlyStopping(monitor='val_prc',mode="max",patience=8,restore_best_weights=True)  
model_5_history = model5.fit(resampled_features,resampled_labels ↪  
↪,batch_size=int(BATCH_SIZE*0.  
↪25),epochs=EPOCHS,callbacks=[lr,early_stopping],validation_data=(X_test ↪  
↪,Y_test),verbose=0)
```

```
[356]: _ = plot_metrics(model_5_history)
```




```
[357]: model_5_results = model5.evaluate(X_test ,Y_test,verbose=0)
```

```
[358]: for name, value in zip(model5.metrics_names, model_5_results):
        print(name, ': {:.3f}'.format(value))
        print()
```

```
loss : 0.699
compile_metrics : 0.735
```

```
[359]: model_5_predictions = np.argmax(model5.predict(X_test,verbose=0), axis=-1)
```

```
[360]: _ = conf_matrix(y_t, model_5_predictions)
```



```
[361]: score_board6 = score_board_df(y_t, model_5_predictions,"RNN-Glove + Upsampling")
score_board = pd.concat([score_board,score_board6],ignore_index=True,
↳sort=False)
score_board
```

```
[361]:
```

	classifier-name	f1_score	precision	recall	accuracy \
0	LinearSVC + Weights	0.785	0.783	0.789	0.789
1	RNN-Word2Vec	0.585	0.647	0.557	0.713
2	RNN-Word2Vec + Weights	0.649	0.657	0.643	0.643

3	RNN-Glove + Weights	0.735	0.764	0.723	0.723
4	RNN-Glove + Weights + other features	0.769	0.787	0.761	0.761
5	RNN-Glove + Upsampling	0.674	0.657	0.701	0.735

Average-method

0	weighted
1	macro
2	weighted
3	weighted
4	weighted
5	macro